

```
In [1]: import pandas as pd
import numpy as np
import matplotlib
```

```
In [2]: #how to import data into python

df=pd.read_csv(r"C:\Users\USER\OneDrive\Documentos\student_habits_performance.csv")
```

```
In [4]: df
```

```
Out[4]:
```

|     | student_id | age | gender | study_hours_per_day | social_media_hours | netflix_hours | part |
|-----|------------|-----|--------|---------------------|--------------------|---------------|------|
| 0   | S1000      | 23  | Female | 0.0                 | 1.2                | 1.1           |      |
| 1   | S1001      | 20  | Female | 6.9                 | 2.8                | 2.3           |      |
| 2   | S1002      | 21  | Male   | 1.4                 | 3.1                | 1.3           |      |
| 3   | S1003      | 23  | Female | 1.0                 | 3.9                | 1.0           |      |
| 4   | S1004      | 19  | Female | 5.0                 | 4.4                | 0.5           |      |
| ... | ...        | ... | ...    | ...                 | ...                | ...           | ...  |
| 995 | S1995      | 21  | Female | 2.6                 | 0.5                | 1.6           |      |
| 996 | S1996      | 17  | Female | 2.9                 | 1.0                | 2.4           |      |
| 997 | S1997      | 20  | Male   | 3.0                 | 2.6                | 1.3           |      |
| 998 | S1998      | 24  | Male   | 5.4                 | 4.1                | 1.1           |      |
| 999 | S1999      | 19  | Female | 4.3                 | 2.9                | 1.9           |      |

1000 rows × 16 columns



```
In [ ]:
```

```
In [ ]:
```

```
In [ ]:
```

```
In [31]: # first step to take after importing your data is to inspect your data


###Data inspection using head,tail,info,describe,columns,index,datatype,shape

#to check the first 5 rows of the data

df.head()
```

Out[31]:

|   | student_id | age | gender | study_hours_per_day | social_media_hours | netflix_hours | part_ti |
|---|------------|-----|--------|---------------------|--------------------|---------------|---------|
| 0 | S1000      | 23  | Female | 0.0                 | 1.2                | 1.1           |         |
| 1 | S1001      | 20  | Female | 6.9                 | 2.8                | 2.3           |         |
| 2 | S1002      | 21  | Male   | 1.4                 | 3.1                | 1.3           |         |
| 3 | S1003      | 23  | Female | 1.0                 | 3.9                | 1.0           |         |
| 4 | S1004      | 19  | Female | 5.0                 | 4.4                | 0.5           |         |




In [32]: *#to check the last 5 rows of the data*

```
df.tail()
```

Out[32]:

|     | student_id | age | gender | study_hours_per_day | social_media_hours | netflix_hours | part |
|-----|------------|-----|--------|---------------------|--------------------|---------------|------|
| 995 | S1995      | 21  | Female | 2.6                 | 0.5                | 1.6           |      |
| 996 | S1996      | 17  | Female | 2.9                 | 1.0                | 2.4           |      |
| 997 | S1997      | 20  | Male   | 3.0                 | 2.6                | 1.3           |      |
| 998 | S1998      | 24  | Male   | 5.4                 | 4.1                | 1.1           |      |
| 999 | S1999      | 19  | Female | 4.3                 | 2.9                | 1.9           |      |



In [33]: *#to check the columns of the data*

```
df.columns
```

Out[33]: Index(['student\_id', 'age', 'gender', 'study\_hours\_per\_day', 'social\_media\_hours', 'netflix\_hours', 'part\_time\_job', 'attendance\_percentage', 'sleep\_hours', 'diet\_quality', 'exercise\_frequency', 'parental\_education\_level', 'internet\_quality', 'mental\_health\_rating', 'extracurricular\_participation', 'exam\_score'], dtype='object')

In [34]: *#to check the shape of the data*

```
df.shape
```

Out[34]: (1000, 16)

In [35]: *#to check the data's information*

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1000 entries, 0 to 999
Data columns (total 16 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   student_id                           1000 non-null   object
1   age                                   1000 non-null   int64
2   gender                               1000 non-null   object
3   study_hours_per_day                  1000 non-null   float64
4   social_media_hours                   1000 non-null   float64
5   netflix_hours                       1000 non-null   float64
6   part_time_job                       1000 non-null   object
7   attendance_percentage                1000 non-null   float64
8   sleep_hours                         1000 non-null   float64
9   diet_quality                        1000 non-null   object
10  exercise_frequency                   1000 non-null   int64
11  parental_education_level             909 non-null    object
12  internet_quality                    1000 non-null   object
13  mental_health_rating                 1000 non-null   int64
14  extracurricular_participation        1000 non-null   object
15  exam_score                          1000 non-null   float64
dtypes: float64(6), int64(3), object(7)
memory usage: 125.1+ KB
```

```
In [36]: #to get describe information of the data

df.describe()
```

Out[36]:

|       | age       | study_hours_per_day | social_media_hours | netflix_hours | attendance_perce |
|-------|-----------|---------------------|--------------------|---------------|------------------|
| count | 1000.0000 | 1000.00000          | 1000.000000        | 1000.000000   | 1000.00          |
| mean  | 20.4980   | 3.55010             | 2.505500           | 1.819700      | 84.13            |
| std   | 2.3081    | 1.46889             | 1.172422           | 1.075118      | 9.39             |
| min   | 17.0000   | 0.00000             | 0.000000           | 0.000000      | 56.00            |
| 25%   | 18.7500   | 2.60000             | 1.700000           | 1.000000      | 78.00            |
| 50%   | 20.0000   | 3.50000             | 2.500000           | 1.800000      | 84.40            |
| 75%   | 23.0000   | 4.50000             | 3.300000           | 2.525000      | 91.00            |
| max   | 24.0000   | 8.30000             | 7.200000           | 5.400000      | 100.00           |

```
In [ ]:
```

```
In [ ]:
```

```
In [ ]:
```

```
In [37]: # Second Step To Take is to clean your data

##check for null,duplicates,change of columns,replace values,drop duplicates,fillna
```

```
#to check for duplicates

df.duplicated()
```

Out[37]:

|     |       |
|-----|-------|
| 0   | False |
| 1   | False |
| 2   | False |
| 3   | False |
| 4   | False |
| ... |       |
| 995 | False |
| 996 | False |
| 997 | False |
| 998 | False |
| 999 | False |

Length: 1000, dtype: bool

In [38]:

```
#to check for duplicates

df[df.duplicated()]
```

Out[38]:

|  | student_id | age | gender | study_hours_per_day | social_media_hours | netflix_hours | part_tin |
|--|------------|-----|--------|---------------------|--------------------|---------------|----------|
|--|------------|-----|--------|---------------------|--------------------|---------------|----------|

In [39]:

```
#to drop duplicates from the data sets

df.drop_duplicates()
```

Out[39]:

|     | student_id | age | gender | study_hours_per_day | social_media_hours | netflix_hours | part |
|-----|------------|-----|--------|---------------------|--------------------|---------------|------|
| 0   | S1000      | 23  | Female | 0.0                 | 1.2                | 1.1           |      |
| 1   | S1001      | 20  | Female | 6.9                 | 2.8                | 2.3           |      |
| 2   | S1002      | 21  | Male   | 1.4                 | 3.1                | 1.3           |      |
| 3   | S1003      | 23  | Female | 1.0                 | 3.9                | 1.0           |      |
| 4   | S1004      | 19  | Female | 5.0                 | 4.4                | 0.5           |      |
| ... | ...        | ... | ...    | ...                 | ...                | ...           | ...  |
| 995 | S1995      | 21  | Female | 2.6                 | 0.5                | 1.6           |      |
| 996 | S1996      | 17  | Female | 2.9                 | 1.0                | 2.4           |      |
| 997 | S1997      | 20  | Male   | 3.0                 | 2.6                | 1.3           |      |
| 998 | S1998      | 24  | Male   | 5.4                 | 4.1                | 1.1           |      |
| 999 | S1999      | 19  | Female | 4.3                 | 2.9                | 1.9           |      |

1000 rows × 16 columns

```
In [40]: #to drop duplicates in a column
df["gender"].drop_duplicates()
```

Out[40]: 0 Female
2 Male
18 Other
Name: gender, dtype: object

```
In [41]: # to check for null values
df.isnull()
```

Out[41]:

|     | student_id | age   | gender | study_hours_per_day | social_media_hours | netflix_hours | par |
|-----|------------|-------|--------|---------------------|--------------------|---------------|-----|
| 0   | False      | False | False  | False               | False              | False         |     |
| 1   | False      | False | False  | False               | False              | False         |     |
| 2   | False      | False | False  | False               | False              | False         |     |
| 3   | False      | False | False  | False               | False              | False         |     |
| 4   | False      | False | False  | False               | False              | False         |     |
| ... | ...        | ...   | ...    | ...                 | ...                | ...           | ... |
| 995 | False      | False | False  | False               | False              | False         |     |
| 996 | False      | False | False  | False               | False              | False         |     |
| 997 | False      | False | False  | False               | False              | False         |     |
| 998 | False      | False | False  | False               | False              | False         |     |
| 999 | False      | False | False  | False               | False              | False         |     |

1000 rows × 16 columns



```
In [42]: #to drop null values
df.dropna()
```

Out[42]:

|     | student_id | age | gender | study_hours_per_day | social_media_hours | netflix_hours | part |
|-----|------------|-----|--------|---------------------|--------------------|---------------|------|
| 0   | S1000      | 23  | Female | 0.0                 | 1.2                | 1.1           |      |
| 1   | S1001      | 20  | Female | 6.9                 | 2.8                | 2.3           |      |
| 2   | S1002      | 21  | Male   | 1.4                 | 3.1                | 1.3           |      |
| 3   | S1003      | 23  | Female | 1.0                 | 3.9                | 1.0           |      |
| 4   | S1004      | 19  | Female | 5.0                 | 4.4                | 0.5           |      |
| ... | ...        | ... | ...    | ...                 | ...                | ...           | ...  |
| 995 | S1995      | 21  | Female | 2.6                 | 0.5                | 1.6           |      |
| 996 | S1996      | 17  | Female | 2.9                 | 1.0                | 2.4           |      |
| 997 | S1997      | 20  | Male   | 3.0                 | 2.6                | 1.3           |      |
| 998 | S1998      | 24  | Male   | 5.4                 | 4.1                | 1.1           |      |
| 999 | S1999      | 19  | Female | 4.3                 | 2.9                | 1.9           |      |

909 rows × 16 columns



In [43]:

```
# to replace null in the data with zero 0
df.fillna(0)
```

Out[43]:

|     | student_id | age | gender | study_hours_per_day | social_media_hours | netflix_hours | part |
|-----|------------|-----|--------|---------------------|--------------------|---------------|------|
| 0   | S1000      | 23  | Female | 0.0                 | 1.2                | 1.1           |      |
| 1   | S1001      | 20  | Female | 6.9                 | 2.8                | 2.3           |      |
| 2   | S1002      | 21  | Male   | 1.4                 | 3.1                | 1.3           |      |
| 3   | S1003      | 23  | Female | 1.0                 | 3.9                | 1.0           |      |
| 4   | S1004      | 19  | Female | 5.0                 | 4.4                | 0.5           |      |
| ... | ...        | ... | ...    | ...                 | ...                | ...           | ...  |
| 995 | S1995      | 21  | Female | 2.6                 | 0.5                | 1.6           |      |
| 996 | S1996      | 17  | Female | 2.9                 | 1.0                | 2.4           |      |
| 997 | S1997      | 20  | Male   | 3.0                 | 2.6                | 1.3           |      |
| 998 | S1998      | 24  | Male   | 5.4                 | 4.1                | 1.1           |      |
| 999 | S1999      | 19  | Female | 4.3                 | 2.9                | 1.9           |      |

1000 rows × 16 columns



In [44]: *# to replace null value in a specific column with zero 0*

```
df["age"].fillna(0)
```

Out[44]:

|     |    |
|-----|----|
| 0   | 23 |
| 1   | 20 |
| 2   | 21 |
| 3   | 23 |
| 4   | 19 |
|     | .. |
| 995 | 21 |
| 996 | 17 |
| 997 | 20 |
| 998 | 24 |
| 999 | 19 |

Name: age, Length: 1000, dtype: int64

In [45]: *#to create a new column and replace a value into the new column from the old column  
#you can then go ahead to drop the replaced column*

```
df["Gender"]=df["gender"].replace("Female","F")
```

In [46]: df

Out[46]:

|     | student_id | age | gender | study_hours_per_day | social_media_hours | netflix_hours | part |
|-----|------------|-----|--------|---------------------|--------------------|---------------|------|
| 0   | S1000      | 23  | Female | 0.0                 | 1.2                | 1.1           |      |
| 1   | S1001      | 20  | Female | 6.9                 | 2.8                | 2.3           |      |
| 2   | S1002      | 21  | Male   | 1.4                 | 3.1                | 1.3           |      |
| 3   | S1003      | 23  | Female | 1.0                 | 3.9                | 1.0           |      |
| 4   | S1004      | 19  | Female | 5.0                 | 4.4                | 0.5           |      |
| ... | ...        | ... | ...    | ...                 | ...                | ...           | ...  |
| 995 | S1995      | 21  | Female | 2.6                 | 0.5                | 1.6           |      |
| 996 | S1996      | 17  | Female | 2.9                 | 1.0                | 2.4           |      |
| 997 | S1997      | 20  | Male   | 3.0                 | 2.6                | 1.3           |      |
| 998 | S1998      | 24  | Male   | 5.4                 | 4.1                | 1.1           |      |
| 999 | S1999      | 19  | Female | 4.3                 | 2.9                | 1.9           |      |

1000 rows × 7 columns



In [ ]:

In [47]: *#to drop a column from the data set  
#after dropping use INPLACE to permanent the change*

```
df.drop(["New_Column"],axis= 1)
```



```

-----
KeyError                                Traceback (most recent call last)
Cell In[47], line 4
      1 #to drop a column from the data set
      2 #after dropping use INPLACE to permanent the change
----> 4 df.drop(["New_Column"],axis= 1)

File ~\anaconda3\Lib\site-packages\pandas\core\frame.py:5581, in DataFrame.drop(self, labels, axis, index, columns, level, inplace, errors)
    5433 def drop(
    5434     self,
    5435     labels: IndexLabel | None = None,
    (... )
    5442     errors: IgnoreRaise = "raise",
    5443 ) -> DataFrame | None:
    5444     """
    5445     Drop specified labels from rows or columns.
    5446
    (... )
    5579         weight  1.0      0.8
    5580     """
-> 5581     return super().drop(
    5582         labels=labels,
    5583         axis=axis,
    5584         index=index,
    5585         columns=columns,
    5586         level=level,
    5587         inplace=inplace,
    5588         errors=errors,
    5589     )

File ~\anaconda3\Lib\site-packages\pandas\core\generic.py:4788, in NDFrame.drop(self, labels, axis, index, columns, level, inplace, errors)
    4786 for axis, labels in axes.items():
    4787     if labels is not None:
-> 4788         obj = obj._drop_axis(labels, axis, level=level, errors=errors)
    4790 if inplace:
    4791     self._update_inplace(obj)

File ~\anaconda3\Lib\site-packages\pandas\core\generic.py:4830, in NDFrame._drop_axis(self, labels, axis, level, errors, only_slice)
    4828     new_axis = axis.drop(labels, level=level, errors=errors)
    4829     else:
-> 4830     new_axis = axis.drop(labels, errors=errors)
    4831     indexer = axis.get_indexer(new_axis)
    4833 # Case for non-unique axis
    4834 else:

File ~\anaconda3\Lib\site-packages\pandas\core\indexes\base.py:7070, in Index.drop(self, labels, errors)
    7068 if mask.any():
    7069     if errors != "ignore":
-> 7070         raise KeyError(f"{labels[mask].tolist()} not found in axis")
    7071     indexer = indexer[~mask]
    7072     return self.delete(indexer)

```

```
KeyError: '['New_Column'] not found in axis"
```

```
In [ ]: #to drop a table permanently you introduce INPLACE=True  
df.drop(["New_Column"],axis=1,inplace=True)
```

```
In [ ]: df
```

```
In [ ]: #to sort values  
df.sort_values("age",ascending=True)
```

```
In [ ]: #to change data type of a particular column if need be  
df["Newtype"]=df["study_hours_per_day"].astype(int)
```

```
In [ ]: df
```

```
In [ ]:
```

```
In [ ]:
```

```
In [ ]:
```

```
In [ ]: #Assignment 1  
#drop the "newcolumn" used to change data type  
df.drop(["Newtype"],axis=1)
```

```
In [ ]: df
```

```
In [ ]: df.drop(["Newtype"],axis=1,inplace=True)
```

```
In [ ]: df
```

```
In [ ]:
```


```
In [ ]:
```

```
In [ ]:
```

```
In [48]: # Third Step To Take is Analysing your Data Based On Giving Measures  
  
## data analysis, statistical analysis  
df.head(3)
```

Out[48]:

|   | student_id | age | gender | study_hours_per_day | social_media_hours | netflix_hours | part_ti |
|---|------------|-----|--------|---------------------|--------------------|---------------|---------|
| 0 | S1000      | 23  | Female | 0.0                 | 1.2                | 1.1           |         |
| 1 | S1001      | 20  | Female | 6.9                 | 2.8                | 2.3           |         |
| 2 | S1002      | 21  | Male   | 1.4                 | 3.1                | 1.3           |         |



In [49]: *# How many students have parttime jobs?*

```
df["part_time_job"].value_counts()
```

Out[49]: part\_time\_job  
No 785  
Yes 215  
Name: count, dtype: int64

In [50]: *# How many female and males are in the data*

```
df["gender"].value_counts()
```

Out[50]: gender  
Female 481  
Male 477  
Other 42  
Name: count, dtype: int64

In [51]: *# what is the maximum study hours per day*

```
df["study_hours_per_day"].max()
```

Out[51]: 8.3

In [52]: *# what is the minimum study hour per day*

```
df["study_hours_per_day"].min()
```

Out[52]: 0.0

In [53]: *# what is the mean age*

```
df["age"].mean()
```

Out[53]: 20.498

In [ ]:

In [ ]:

In [ ]:

```
In [54]: #ASSIGNMENT 2
#SHOW VALUE COUNT OF DIET QUALITY FOR ALL STUDENTS
#AVERAGE ATTENDANCE PERCENTAGE FOR STUDENTS
#CORRELATION ANALYSIS OF NUMERICS

df.head()
```

```
Out[54]:
```

|   | student_id | age | gender | study_hours_per_day | social_media_hours | netflix_hours | part_ti |
|---|------------|-----|--------|---------------------|--------------------|---------------|---------|
| 0 | S1000      | 23  | Female | 0.0                 | 1.2                | 1.1           |         |
| 1 | S1001      | 20  | Female | 6.9                 | 2.8                | 2.3           |         |
| 2 | S1002      | 21  | Male   | 1.4                 | 3.1                | 1.3           |         |
| 3 | S1003      | 23  | Female | 1.0                 | 3.9                | 1.0           |         |
| 4 | S1004      | 19  | Female | 5.0                 | 4.4                | 0.5           |         |

```
In [55]: #Value Count Of Diet For All Students

df[["diet_quality"]].value_counts()
```

```
Out[55]: diet_quality
Fair      437
Good      378
Poor      185
Name: count, dtype: int64
```

```
In [56]: # Average attendance percentage of all students

df[["attendance_percentage"]].mean()
```

```
Out[56]: attendance_percentage    84.1317
dtype: float64
```

```
In [57]: # Average attendance percentage of all gender

df.groupby(["gender"])[["attendance_percentage"]].mean()
```

```
Out[57]:
```

| attendance_percentage |           |
|-----------------------|-----------|
| gender                |           |
| Female                | 84.371518 |
| Male                  | 83.894549 |
| Other                 | 84.078571 |

```
In [58]: # Average attendance percentage of all students

df.groupby(["gender", "age"])[["attendance_percentage"]].mean()
```

Out[58]:

| attendance_percentage |     |           |
|-----------------------|-----|-----------|
| gender                | age |           |
| Female                | 17  | 85.758621 |
|                       | 18  | 84.340678 |
|                       | 19  | 84.922951 |
|                       | 20  | 83.929730 |
|                       | 21  | 84.360317 |
|                       | 22  | 83.068333 |
|                       | 23  | 84.476471 |
|                       | 24  | 84.261818 |
| Male                  | 17  | 83.256522 |
|                       | 18  | 86.331481 |
|                       | 19  | 82.710204 |
|                       | 20  | 83.435385 |
|                       | 21  | 83.023636 |
|                       | 22  | 83.488235 |
|                       | 23  | 85.308065 |
|                       | 24  | 83.634722 |
| Other                 | 17  | 88.666667 |
|                       | 18  | 89.575000 |
|                       | 19  | 82.766667 |
|                       | 20  | 80.871429 |
|                       | 21  | 78.228571 |
|                       | 22  | 90.600000 |
|                       | 23  | 87.300000 |
|                       | 24  | 82.000000 |

In [59]:

```
#CORRELATION ANALYSIS OF NUMERICS
## To reveal all columns that contain numerical data

df.select_dtypes(include=['int64', 'float64']).corr()
```

Out[59]:

|                       | age       | study_hours_per_day | social_media_hours | netflix_hours | attendance_percentage | sleep_hours | exercise_frequency | mental_health_rating | exam_score |
|-----------------------|-----------|---------------------|--------------------|---------------|-----------------------|-------------|--------------------|----------------------|------------|
| age                   | 1.000000  | 0.003971            | -0.009151          | -0.001174     | -0.026055             | 0.037482    | -0.003836          | -0.045101            | -0.008907  |
| study_hours_per_day   | 0.003971  | 1.000000            | 0.020282           | -0.031158     | 0.026264              | -0.027757   | -0.028701          | -0.003768            | 0.825419   |
| social_media_hours    | -0.009151 | 0.020282            | 1.000000           | 0.011477      | 0.040479              | 0.018236    | -0.037319          | 0.001496             | -0.166733  |
| netflix_hours         | -0.001174 | -0.031158           | 0.011477           | 1.000000      | -0.002092             | -0.000935   | -0.006448          | 0.008034             | -0.171779  |
| attendance_percentage | -0.026055 | 0.026264            | 0.040479           | -0.002092     |                       |             |                    |                      |            |
| sleep_hours           | 0.037482  | -0.027757           | 0.018236           | -0.000935     |                       |             |                    |                      |            |
| exercise_frequency    | -0.003836 | -0.028701           | -0.037319          | -0.006448     |                       |             |                    |                      |            |
| mental_health_rating  | -0.045101 | -0.003768           | 0.001496           | 0.008034      |                       |             |                    |                      |            |
| exam_score            | -0.008907 | 0.825419            | -0.166733          | -0.171779     |                       |             |                    |                      |            |

In [ ]:

In [60]:

df.head(2)

Out[60]:

|   | student_id | age | gender | study_hours_per_day | social_media_hours | netflix_hours | part_time_job |
|---|------------|-----|--------|---------------------|--------------------|---------------|---------------|
| 0 | S1000      | 23  | Female | 0.0                 | 1.2                | 1.1           |               |
| 1 | S1001      | 20  | Female | 6.9                 | 2.8                | 2.3           |               |

In [61]:

```
# Sort
## total nextflix hours of all students sort by hours

df.groupby("gender")[["netflix_hours"]].sum().sort_values("netflix_hours",ascending=True)
```

Out[61]:

| netflix_hours |       |
|---------------|-------|
| gender        |       |
| Male          | 868.5 |
| Female        | 864.8 |
| Other         | 86.4  |

In [62]:

```
# check exam score for all gender and parttime job

df.groupby(["gender","part_time_job"])[["exam_score"]].sum()
```

Out[62]:

| exam_score |               |         |
|------------|---------------|---------|
| gender     | part_time_job |         |
| Female     | No            | 26806.6 |
|            | Yes           | 6739.0  |
| Male       | No            | 25692.7 |
|            | Yes           | 7396.0  |
| Other      | No            | 2322.4  |
|            | Yes           | 644.8   |

In [63]: *# check exam score for all gender*

```
df.groupby(["gender", "age"])[["exam_score"]].sum()
```

Out[63]:

| exam_score |     |        |
|------------|-----|--------|
| gender     | age |        |
| Female     | 17  | 4140.5 |
|            | 18  | 4279.9 |
|            | 19  | 4203.5 |
|            | 20  | 5273.1 |
|            | 21  | 4329.2 |
|            | 22  | 4142.7 |
|            | 23  | 3505.1 |
|            | 24  | 3671.6 |
| Male       | 17  | 4829.9 |
|            | 18  | 3706.4 |
|            | 19  | 3325.3 |
|            | 20  | 4471.2 |
|            | 21  | 3639.9 |
|            | 22  | 3530.2 |
|            | 23  | 4444.4 |
|            | 24  | 5141.4 |
| Other      | 17  | 385.2  |
|            | 18  | 313.6  |
|            | 19  | 188.4  |
|            | 20  | 464.0  |
|            | 21  | 473.9  |
|            | 22  | 135.7  |
|            | 23  | 487.8  |
|            | 24  | 518.6  |

In [64]: `df.head(2)`



Out[64]:

|   | student_id | age | gender | study_hours_per_day | social_media_hours | netflix_hours | part_ti |
|---|------------|-----|--------|---------------------|--------------------|---------------|---------|
| 0 | S1000      | 23  | Female | 0.0                 | 1.2                | 1.1           |         |
| 1 | S1001      | 20  | Female | 6.9                 | 2.8                | 2.3           |         |

In [65]:

```
df.pivot_table(index=["gender"],columns=["part_time_job"],values=["exam_score"],agg
```

Out[65]:

| sum           |         |        |
|---------------|---------|--------|
| exam_score    |         |        |
| part_time_job | No      | Yes    |
| gender        |         |        |
| Female        | 26806.6 | 6739.0 |
| Male          | 25692.7 | 7396.0 |
| Other         | 2322.4  | 644.8  |

In [ ]:

In [ ]:

In [ ]:

In [66]:

```
#ASSIGNMENT 3
#CALCULATE THE AVERAGE SLEEP HOURS FOR THOSE ON DIET QUALITY
#CALCULATE THE TOTAL STUDENT UNDER DIET QUALITY
#WHICH EDUCATION LEVEL HAD THE HIGHEST AGE OF ALL STUDENT

df.head(2)
```

Out[66]:

|   | student_id | age | gender | study_hours_per_day | social_media_hours | netflix_hours | part_ti |
|---|------------|-----|--------|---------------------|--------------------|---------------|---------|
| 0 | S1000      | 23  | Female | 0.0                 | 1.2                | 1.1           |         |
| 1 | S1001      | 20  | Female | 6.9                 | 2.8                | 2.3           |         |

In [67]:

```
# Average sleep hours for those on diet
df.groupby("diet_quality")[["sleep_hours"]].mean().sort_values("sleep_hours",ascend
```

Out[67]: **sleep\_hours**

| <b>diet_quality</b> |          |
|---------------------|----------|
| <b>Poor</b>         | 6.559459 |
| <b>Fair</b>         | 6.465217 |
| <b>Good</b>         | 6.432011 |

```
In [68]: # Total student under diet quality
df[["gender", "diet_quality"]].value_counts()
```

```
Out[68]: gender  diet_quality
Female  Fair          210
Male    Fair          209
         Good          185
Female  Good          179
         Poor           92
Male    Poor           83
Other   Fair           18
         Good           14
         Poor           10
Name: count, dtype: int64
```

```
In [33]: # which education level had the highest age for all students
df.groupby("parental_education_level")[["age"]].max()
```

Out[33]: **age**

| <b>parental_education_level</b> |    |
|---------------------------------|----|
| <b>Bachelor</b>                 | 24 |
| <b>High School</b>              | 24 |
| <b>Master</b>                   | 24 |

```
In [35]: # which gender had the maximum age in parental education level
df.groupby(["gender", "parental_education_level"])[["age"]].max()
```

Out[35]:

|        |                          | age |
|--------|--------------------------|-----|
| gender | parental_education_level |     |
| Female | Bachelor                 | 24  |
|        | High School              | 24  |
|        | Master                   | 24  |
| Male   | Bachelor                 | 24  |
|        | High School              | 24  |
|        | Master                   | 24  |
| Other  | Bachelor                 | 23  |
|        | High School              | 24  |
|        | Master                   | 24  |

```
In [40]: df.groupby(["gender", "parental_education_level"])[["study_hours_per_day"]].max()
```

Out[40]:

|        |                          | study_hours_per_day |
|--------|--------------------------|---------------------|
| gender | parental_education_level |                     |
| Female | Bachelor                 | 7.5                 |
|        | High School              | 7.6                 |
|        | Master                   | 6.7                 |
| Male   | Bachelor                 | 6.8                 |
|        | High School              | 8.3                 |
|        | Master                   | 8.2                 |
| Other  | Bachelor                 | 4.7                 |
|        | High School              | 5.1                 |
|        | Master                   | 5.6                 |

```
In [41]: # describe age for the parental educational level

df.groupby("parental_education_level")["age"].describe()
```

Out[41]:

|                          | count | mean      | std      | min  | 25%  | 50%  | 75%  | max  |
|--------------------------|-------|-----------|----------|------|------|------|------|------|
| parental_education_level |       |           |          |      |      |      |      |      |
| Bachelor                 | 350.0 | 20.600000 | 2.323420 | 17.0 | 19.0 | 20.0 | 23.0 | 24.0 |
| High School              | 392.0 | 20.364796 | 2.256468 | 17.0 | 18.0 | 20.0 | 22.0 | 24.0 |
| Master                   | 167.0 | 20.473054 | 2.366468 | 17.0 | 18.0 | 20.0 | 23.0 | 24.0 |

In [ ]:

In [ ]:

In [ ]:

In [ ]:

df.head(4)

In [16]:

```
# To apply upper case to a specific columnn

df["applylamnda"]=df["gender"].apply(lambda x:x.upper())
```

In [23]:

df

Out[23]:

|     | student_id | age | gender | study_hours_per_day | social_media_hours | netflix_hours | part |
|-----|------------|-----|--------|---------------------|--------------------|---------------|------|
| 0   | S1000      | 23  | Female | 0.0                 | 1.2                | 1.1           |      |
| 1   | S1001      | 20  | Female | 6.9                 | 2.8                | 2.3           |      |
| 2   | S1002      | 21  | Male   | 1.4                 | 3.1                | 1.3           |      |
| 3   | S1003      | 23  | Female | 1.0                 | 3.9                | 1.0           |      |
| 4   | S1004      | 19  | Female | 5.0                 | 4.4                | 0.5           |      |
| ... | ...        | ... | ...    | ...                 | ...                | ...           |      |
| 995 | S1995      | 21  | Female | 2.6                 | 0.5                | 1.6           |      |
| 996 | S1996      | 17  | Female | 2.9                 | 1.0                | 2.4           |      |
| 997 | S1997      | 20  | Male   | 3.0                 | 2.6                | 1.3           |      |
| 998 | S1998      | 24  | Male   | 5.4                 | 4.1                | 1.1           |      |
| 999 | S1999      | 19  | Female | 4.3                 | 2.9                | 1.9           |      |

1000 rows × 7 columns

In [14]:

```
# Apply a transformation to the mental_health_rating column by multiplying each val
# The apply() method executes the lambda function on each element of the column
# The result is stored in a new column called "Numericapply"
```

```
df["Numericapply"] = df["mental_health_rating"].apply(lambda x: x*2)
```

In [15]: `df.head()`

Out[15]:

|  | student_id | age | gender | study_hours_per_day | social_media_hours | netflix_hours | part_ti |
|--|------------|-----|--------|---------------------|--------------------|---------------|---------|
|--|------------|-----|--------|---------------------|--------------------|---------------|---------|

|   |       |    |        |     |     |     |  |
|---|-------|----|--------|-----|-----|-----|--|
| 0 | S1000 | 23 | Female | 0.0 | 1.2 | 1.1 |  |
| 1 | S1001 | 20 | Female | 6.9 | 2.8 | 2.3 |  |
| 2 | S1002 | 21 | Male   | 1.4 | 3.1 | 1.3 |  |
| 3 | S1003 | 23 | Female | 1.0 | 3.9 | 1.0 |  |
| 4 | S1004 | 19 | Female | 5.0 | 4.4 | 0.5 |  |



In [9]: *# Group the dataframe by 'parental\_education\_level' column  
# Then count the occurrences of each gender within each education level group*

```
df.groupby("parental_education_level")["gender"].count()
```

Out[9]:

|                          |     |
|--------------------------|-----|
| parental_education_level |     |
| Bachelor                 | 350 |
| High School              | 392 |
| Master                   | 167 |

Name: gender, dtype: int64

In [73]: `df.head()`

Out[73]:

|  | student_id | age | gender | study_hours_per_day | social_media_hours | netflix_hours | part_ti |
|--|------------|-----|--------|---------------------|--------------------|---------------|---------|
|--|------------|-----|--------|---------------------|--------------------|---------------|---------|

|   |       |    |        |     |     |     |  |
|---|-------|----|--------|-----|-----|-----|--|
| 0 | S1000 | 23 | Female | 0.0 | 1.2 | 1.1 |  |
| 1 | S1001 | 20 | Female | 6.9 | 2.8 | 2.3 |  |
| 2 | S1002 | 21 | Male   | 1.4 | 3.1 | 1.3 |  |
| 3 | S1003 | 23 | Female | 1.0 | 3.9 | 1.0 |  |
| 4 | S1004 | 19 | Female | 5.0 | 4.4 | 0.5 |  |



In [8]: *# Calculate the mean of the 'age' column and count the number of entries in the 'ge  
# This aggregation returns a new DataFrame with one row containing these summary st*

```
df.agg({"age": "mean", "gender": "count"})
```

Out[8]:

|        |          |
|--------|----------|
| age    | 20.498   |
| gender | 1000.000 |

dtype: float64

In [7]: *# Add 10 to each value in the 'age' column using transform  
# transform applies the function to each element and returns a Series with the same*

```
df["age"].transform(lambda x:x +10)
```

Out[7]:

|     |    |
|-----|----|
| 0   | 33 |
| 1   | 30 |
| 2   | 31 |
| 3   | 33 |
| 4   | 29 |
|     | .. |
| 995 | 31 |
| 996 | 27 |
| 997 | 30 |
| 998 | 34 |
| 999 | 29 |

Name: age, Length: 1000, dtype: int64

```
In [6]: #Correlation coefficient between 'age',and 'exercise_frequency' columns

df["age"].corr(df["exercise_frequency"])
```

Out[6]: -0.0038362358530908297

```
In [32]: # Convert all values in the 'diet_quality' column to lowercase for consistency in d

df["diet_quality"]=df["diet_quality"].apply(lambda x:x.lower())
```

```
In [44]: df
```

Out[44]:

|     | student_id | age | gender | study_hours_per_day | social_media_hours | netflix_hours | part |
|-----|------------|-----|--------|---------------------|--------------------|---------------|------|
| 0   | S1000      | 23  | Female | 0.0                 | 1.2                | 1.1           |      |
| 1   | S1001      | 20  | Female | 6.9                 | 2.8                | 2.3           |      |
| 2   | S1002      | 21  | Male   | 1.4                 | 3.1                | 1.3           |      |
| 3   | S1003      | 23  | Female | 1.0                 | 3.9                | 1.0           |      |
| 4   | S1004      | 19  | Female | 5.0                 | 4.4                | 0.5           |      |
| ... | ...        | ... | ...    | ...                 | ...                | ...           | ...  |
| 995 | S1995      | 21  | Female | 2.6                 | 0.5                | 1.6           |      |
| 996 | S1996      | 17  | Female | 2.9                 | 1.0                | 2.4           |      |
| 997 | S1997      | 20  | Male   | 3.0                 | 2.6                | 1.3           |      |
| 998 | S1998      | 24  | Male   | 5.4                 | 4.1                | 1.1           |      |
| 999 | S1999      | 19  | Female | 4.3                 | 2.9                | 1.9           |      |

1000 rows × 8 columns



```
In [ ]:
```

